

Introduction to Modern CPU architecture

Luca Tornatore, I.N.A.F.

2024 INAF Course on HPC




June, 24th - July 5th, IRA, Bologna



Modern
Arch.

Moore's law, is it over ?

40 Years of Microprocessor Trend Data



« every 18 months
the density of
transistors will
double *at the same
manufacturing cost* »

It is still in place,
it seems


Data collected up to 2010 by M.Horowitz,
F. Labonte, O. Schacham, K. Olukotun,
L.Hammond and C. Batlen.
New data up to 2015 collected by
2010-2015 by K. Rupp.



Modern
Arch.

Moore's law, is it over ?

40 Years of Microprocessor Trend Data



« every 18 months
the density of
transistors will
double *at the same
manufacturing cost* »

It is still in place,
it seems,

with a
fundamental
factor

Data collected up to 2010 by M.Horowitz,
F. Labonte, O. Schacham, K. Olukotun,
L.Hammond and C. Batlen.
New data up to 2015 collected by
2010-2015 by K. Rupp.



Modern Arch.


Why there is no more “free lunch”?

« free lunch » is the fact that you do not act on your code; you just wait 18 months and it will run faster

Power consumption:

$$P \propto C \cdot V^2 \cdot f$$

\downarrow \downarrow \downarrow
 $\times \sim 30$ $5V \rightarrow 1V$ $\times \sim 1000$





| The transition from the “free lunch”

Until half of the '00s, engineers succeeded in gaining performance by essentially 3 ways:

1. Increasing **clock** speed
2. Optimizing **execution**
3. Enlarging/improving **cache**



| The transition from the “free lunch”

Until half of the '00s, engineers succeeded in gaining performance by essentially 3 ways:

1. Increasing **clock** speed →
2. Optimizing **execution**
3. Enlarging/improving **cache**

You get more cycles per unit time;
more or less that means doing the
same bunch of instructions faster



| The transition from the “free lunch”

Until half of the '00s, engineers
gaining performance by essentially

1. Increasing **clock** speed
2. Optimizing **execution** →
3. Enlarging/improving **cache**

- More powerful instructions
- Pipelining
- Branch predictions
- Out-of-order execution
- Vector instructions

Under enormous pressure, CPUs
manufacturers risked (and did) to
break the semantic of your code.
Or introduce horrible bugs..
(have you heard about **Meltdown** and
Spectre ?)



| The transition from the “free lunch”

Until half of the '00s, engineers succeeded in gaining performance by essentially 3 ways:

1. Increasing **clock** speed

2. Optimizing **execution**

3. Enlarging/improving **cache** →

More on that later..



| Why there is no more “free lunch”?

Applications no longer get more performance for free without significant re-design, since $\gtrsim 15$ years

Since 15 years, the gain in performance is essentially due to fundamentally different factors:

1. Multi-core + Multi-threads
2. Enlarging/improving cache
3. Hyperthreading (*smaller contribution*)



| Why there is no more “free lunch”?

For instance:

2 Cores at 3GHz are
basically 1 Core at 6GHz.. ?

False

- ✗ Cores coordination for cache-coherence
- ✗ Threads coordination
- ✗ Memory access
- ✗ Increased algorithmic complexity

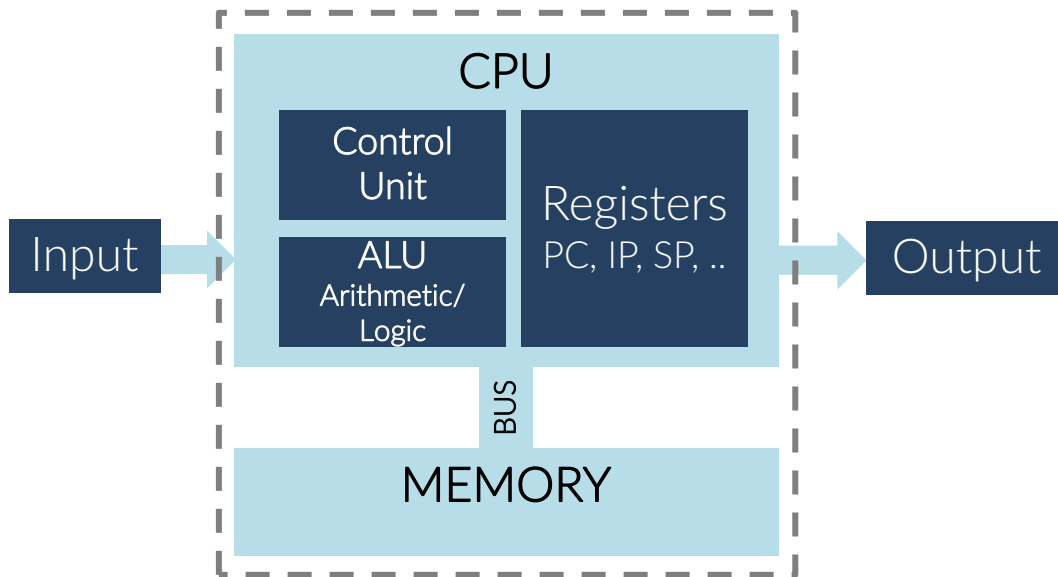
Since 15 years, the gain in performance is essentially due to
fundamentally different factors:

1. Multi-core + **Multi-threads**
2. Enlarging/improving **cache**
3. Hyperthreading (*smaller contribution*)



Once upon a time..

.. The computer architecture was much simpler than today. In the lecture “The Free lunch is over”, we have seen the reasons why it was so – or, better, why nowadays it is much more complex.




In the Von Neumann architecture

- there is only 1 processing unit
- 1 instructions is executed at a time
- memory is “flat”:
 - access on any location has always the same cost
 - access to memory has the same cost than op execution



While today...



NOT an extreme example: 6-cores Intel Xeon e5600

(*) “cost” is in terms of the CPU-cycles currency

- there are *many* processing units
- many instructions can be executed at a time
- many data can be processed at a time
- “instructions” are internally broken down in many simpler μ ops that are pipelined
 - different instructions could have quite different cost^(*)
- memory is strongly *not* “flat”:
 - there is a strong memory hierarchy
 - access memory can have very different costs^(*) depending on the location
 - accessing RAM is way more costly than performing a μ op or reading an internal register

In the next slides we'll go through all this features in chronological order, as they developed in time




Some of the steps in between

- The memory hierarchy
- Superscalability and out-of-order execution
- Pipelining
- Vector capabilities



Early 90s: CPU get faster than memory



Processor	Alpha 21164	
Machine	AlphaServer 8200	
Clock Rate	300 MHz	
Memory Performance	Latency	Bandwidth
I Cache (8KB on chip)	6.7 ns (2 clocks)	4800 MB/sec
D Cache (8KB on chip)	6.7 ns (2 clocks)	4800 MB/sec
L2 Cache (96KB on chip)	20 ns (6 clocks)	4800 MB/sec
L3 Cache (4MB off chip)	26 ns (8 clocks)	960 MB/sec
Main Memory Subsystem	253 ns (76 clocks)	1200 MB/sec
Single DRAM component	$\approx 60\text{ns}$ (18 clocks)	$\approx 30\text{--}100\text{ MB/sec}$

Taken from a 1997 paper

The CPU had to spend more time waiting for data coming from RAM than executing ops. That is part of the so called "memory wall".
What is the solution ?



A note about today

You certainly know that DDRAM memory caught up: frequency has increased in the last year up to $\geq 3\text{-}4\text{GHz}$

However, that alleviate the problem only partially




Figure 2.4: Increasing discrepancy between main memory access latency and CPU cycle time, for last 25 years. Recently the downtrend in CPU cycle time decreased to 2%, while single processor performance improves at 3.5% per year [52]. The decrement in memory access latency is less than 1%.

Figure taken from:
"Memory Bandwidth and Latency in HPC: System Requirements and Performance Impact"

Ph.D dissertation thesis, 2019, Dep.t of Computing Architectures, UPC




The cache memory

The solution is to equip your CPU with a *faster* memory, that, by the way, is also way more expensive and way more energy consuming.

Furthermore, to be faster it ought to be *extremely closer*. All in all, the new memory that will be called **cache**, will be much *smaller* than RAM.

The cache itself has a hierarchy:

- *Level-I* is inside each core.
- *Level-II* is also inside the core, or may be shared by more cores.
- *Level-III* is inside the CPU, shared by many cores.



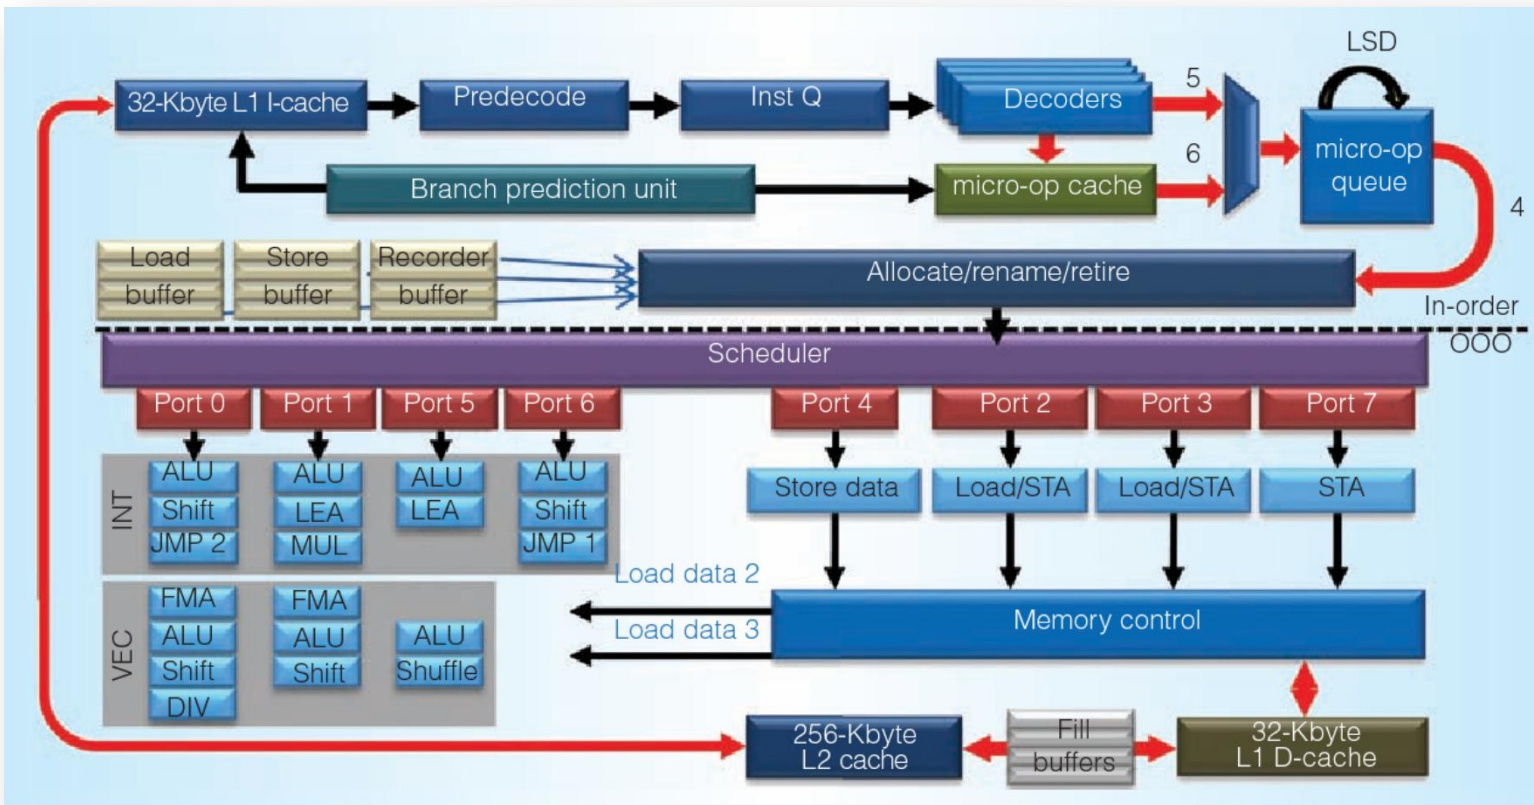


Some of the steps in between

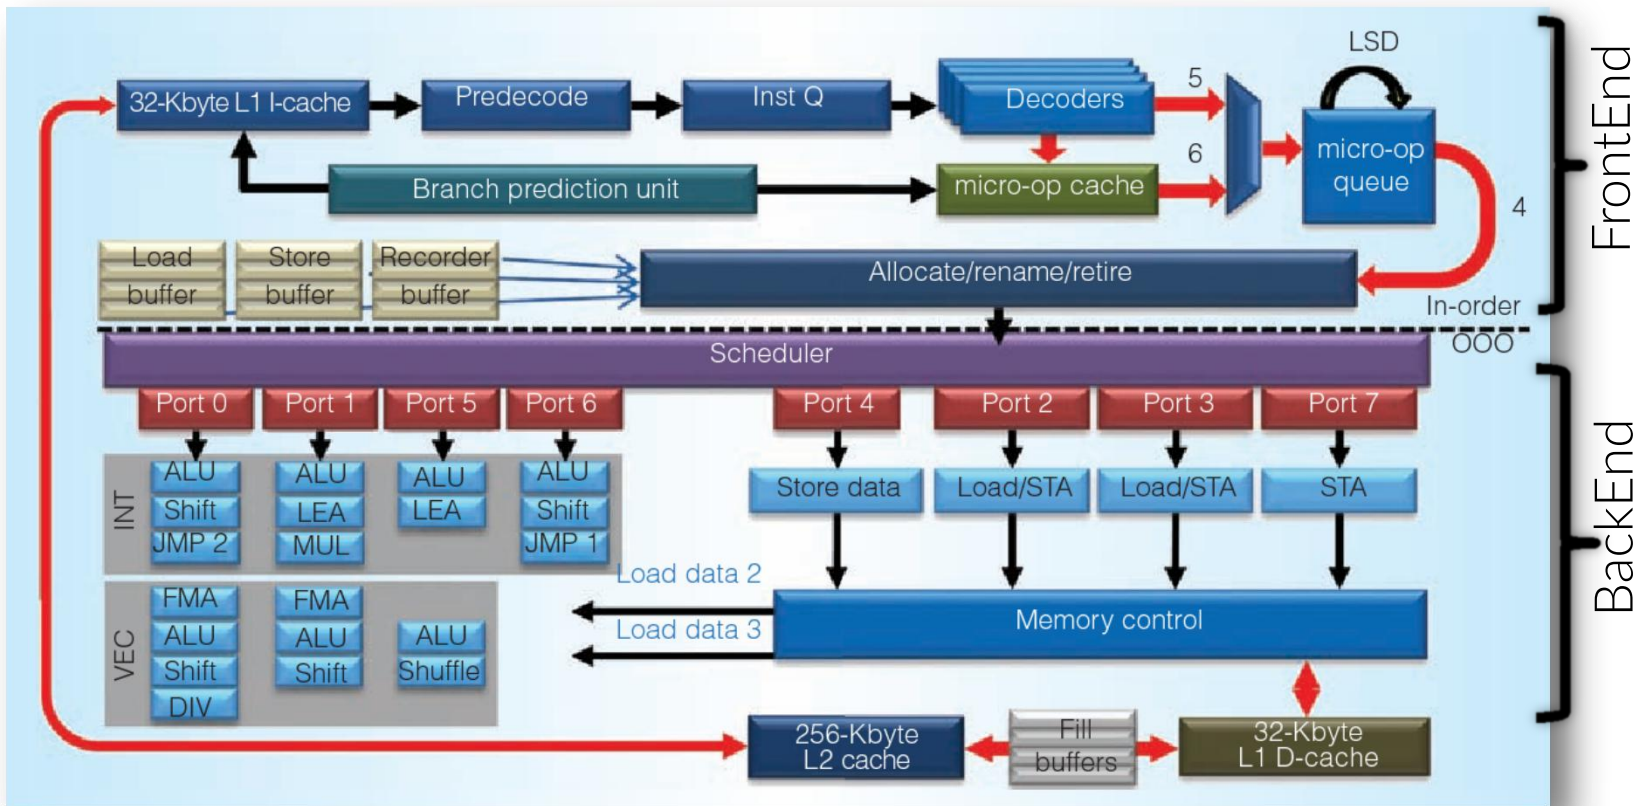
- The memory hierarchy
- Superscalability and out-of-order execution
- Pipelining
- Vector capabilities



Mid 90s: superscalar and out-of-order CPUs

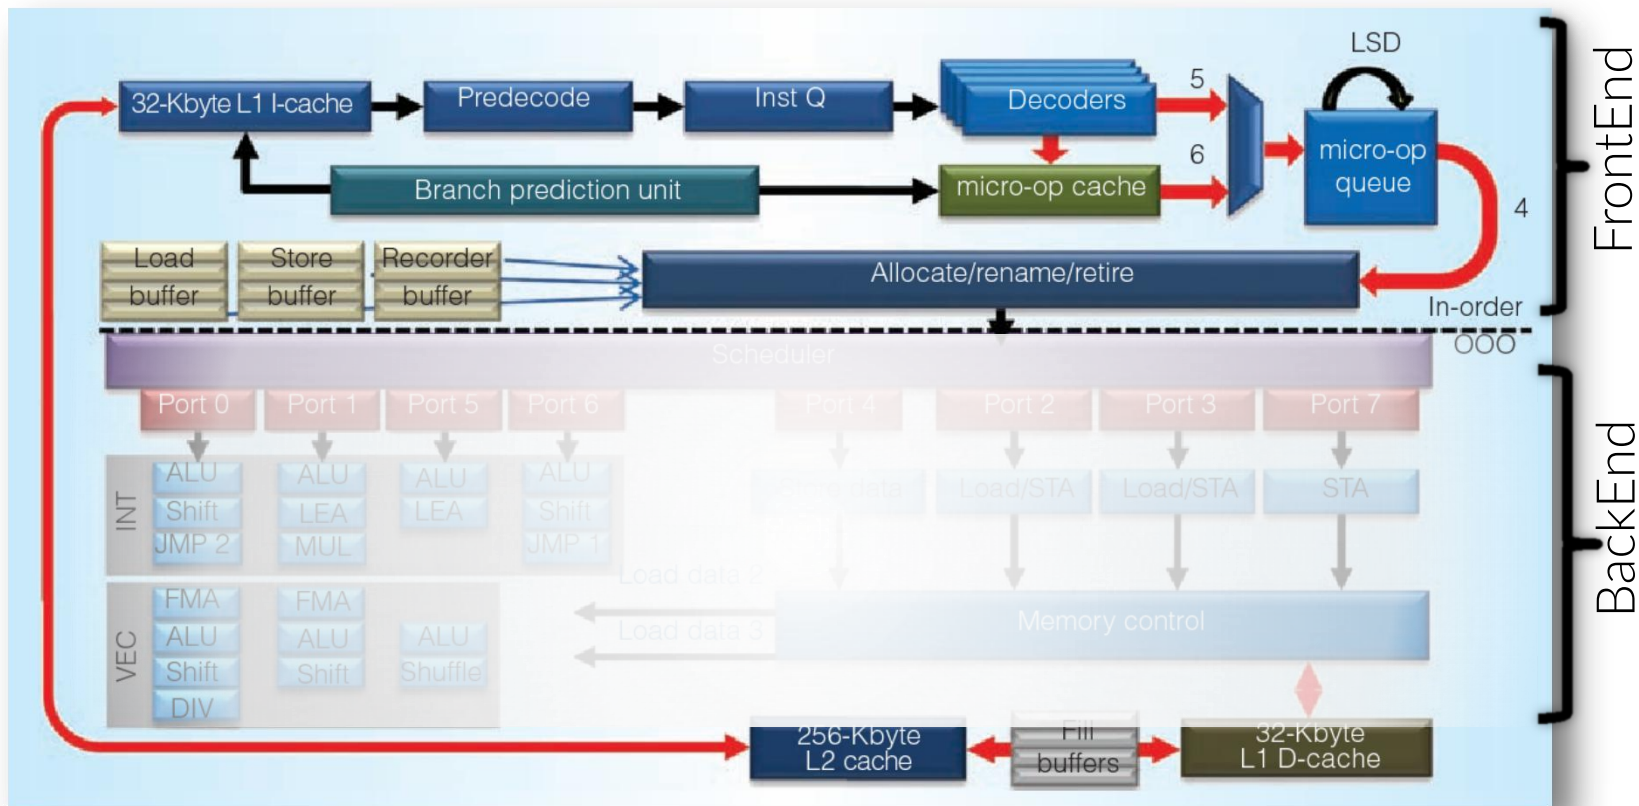


More than 1 *port* is available to execute CPU instructions, although different units have different specializations (ALU, LEA, SHIFT, FMA, ...) : that is **superscalar capacity**, i.e. the capacity of executing more than 1 instructions per cycle.



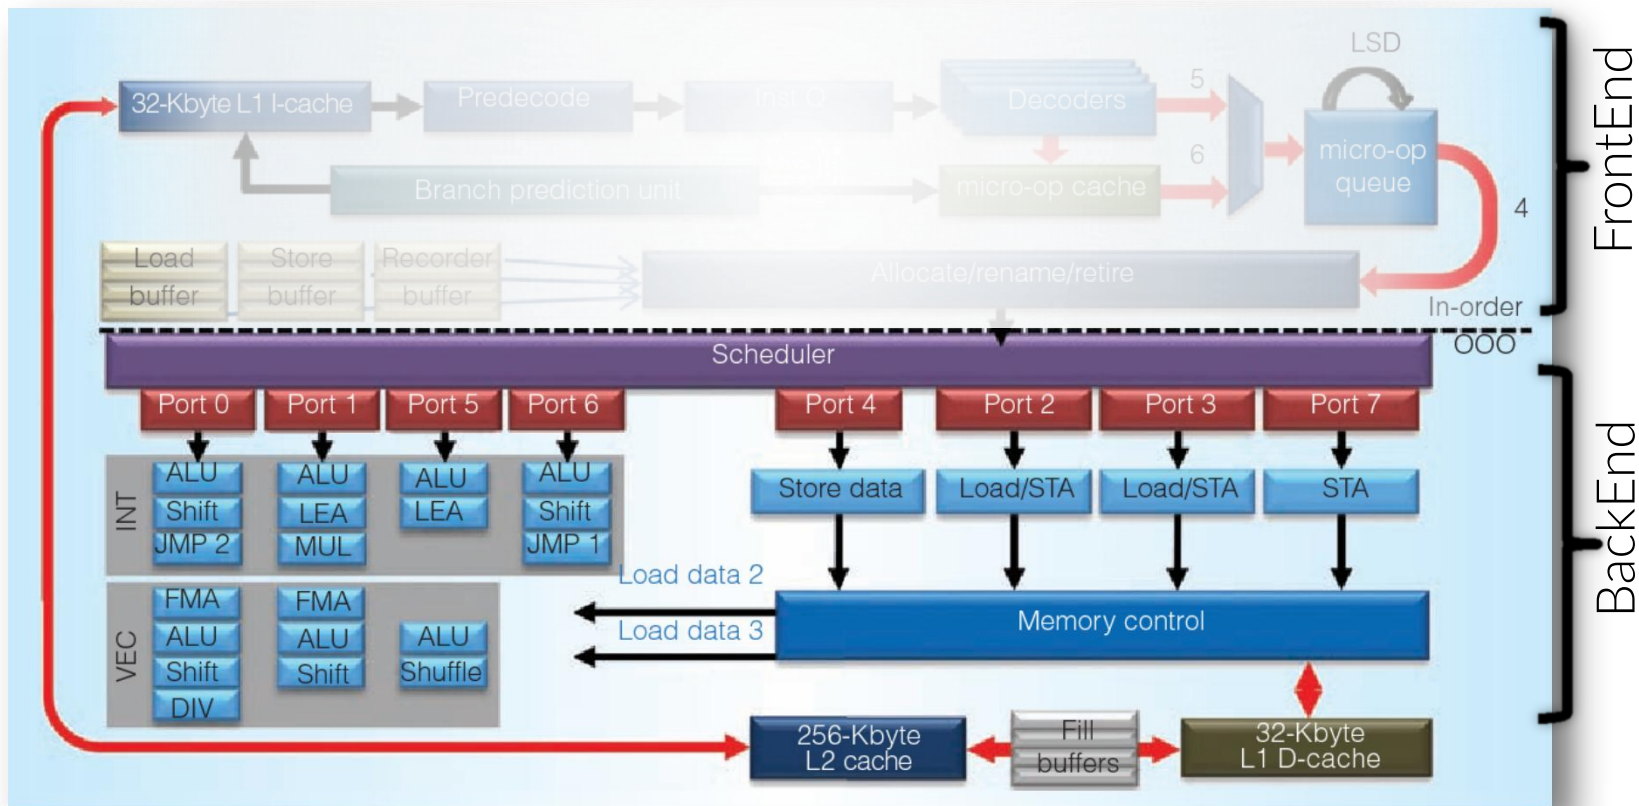
6th generation
SkyLake
micro-arch.

The *Front-End* basically fetches instructions and the data they operate on from instruction and data caches, decodes instructions, predicts branches and dispatches the instructions to different ports



6th generation
SkyLake
micro-arch.

The *Back-End* is responsible for the actual instructions execution and for the back-writing of results in memory locations. It is responsible also for orchestrating **out-of-order** ops execution depending on their instructions/data dependencies.



6th generation
SkyLake
micro-arch.



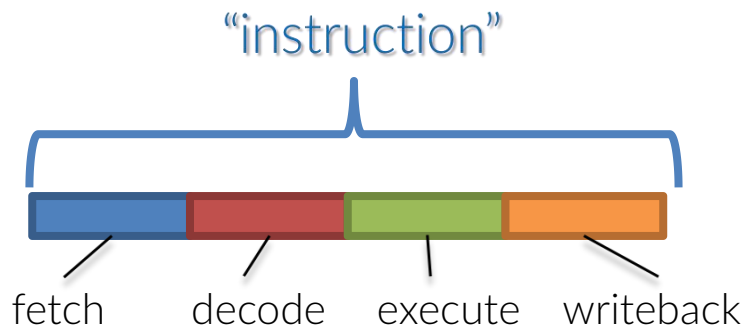
Some of the steps in between

- The memory hierarchy
- Superscalability and out-of-order execution
- Pipelining
- Vector capabilities




It would be obvious to think that an “instruction” is a kind of *atomic* operation that the CPU perform as a whole. Indeed that was true until the mid of the 80s.

If you think carefully about it, it is easy to understand that actually an “instruction” involves at least the following *independent* steps:



1. *Fetching*
it must be recalled from memory/lcache
2. *Decoding:*
it must be “understood and interpreted”
3. *Execution*
4. *Writeback :*
the result must be accounted in memory ()





If all the four stages take $\sim 400\text{ps}$, we then would obtain a **throughput** of 2.5GIPS (giga-instructions per second).

400ps is also the total time required to get a result from an instruction, and so it is the **latency** of the instruction.

However, if we were able to “detach” the four stages, we could organize things differently, like in a car-building chain, or even at the mensa of the university.



Note: all the timing estimates are hypothetical for the purpose of discussing the concepts



If many independent logical units exist to perform each step, they could operate subsequently on **different instructions**:


If the stage delays are not uniform, the throughput is limited by the latency $F + (D + E) - (F + D) = E \sim 220\text{ps}$, which means we have a throughput of **$\sim 4.5\text{GIPS}$** just because of logic units separation.



Therefore, introducing the instructions pipelining, we can increase the **throughput** of our system by a large factor.

However, the efficiency of the pipelines is limited by its longest stage: the better option would be to have all equal stages, for instance further subdividing each stage – especially the most demanding ones (*).

increase due to
additional logics



Now the throughput of our system has increased to 1 instruction retired every 50ps, i.e. **20GIPS**

(*) this is called *superpipelining*: modern CPUs may have 10-20 stages per pipeline.



Pipelines are about throughput

Pipelining is then about increasing the throughput of a system, and as we have seen it could be really effective.

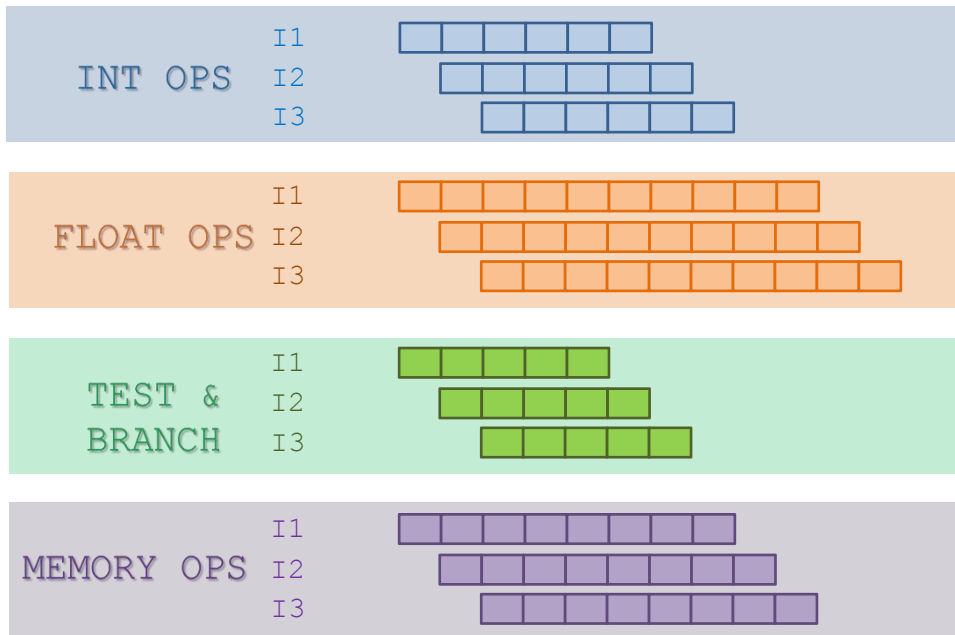

However, there is more that can be done working on the **execution stage** that, of course, encompasses a large number of different *functional units*, performing a **different and independent tasks**.

With an enhancement of the decode/dispatch stage, we can address multiple pipelines that can be active at the same time:

DISCLAIMER: *of course things are more complicated than depicted here, and we do not get 10^{10} GIPS.*



Multiple pipelines





Some of the steps in between


- The memory hierarchy
- Superscalability and out-of-order execution
- Pipelining
- Vector capabilities




Late 90s: vector capabilities become mainstream

https://software.intel.com/sites/default/files/m/d/4/1/d/8/Intro_to_Intel_AVX.pdf

SIMD Mode



Scalar Mode



Vector registers are large special registers in the CPU that can be considered subdivided in smaller independent chunks over which the same operation can be performed:

SIMD: Single Instruction
Multiple Data



Vector registers size

SSE Data Types (16 XMM Registers)

__m128	Float	Float	Float	Float	4x 32-bit float												
__m128d	Double		Double		2x 64-bit double												
__m128i	B	B	B	B	B	B	B	B	B	B	B	B	B	B	B	16x 8-bit byte	
__m128i	short	short	short	short	short	short	short	short	short	short	short	short	short	short	short	8x 16-bit short	
__m128i	int		int		int		int		int		int		int		4x 32bit integer		
__m128i	long long				long long				long long				long long				2x 64bit long
__m128i	doublequadword																1x 128-bit quad


AVX Data Types (16 YMM Registers)

__mm256	Float	Float	Float	Float	Float	Float	Float	Float	8x 32-bit float
__mm256d	Double		Double		Double		Double		4x 64-bit double
__mm256i	256-bit Integer registers. It behaves similarly to __m128i. Out of scope in AVX, useful on AVX2								




Vector registers size


a



b

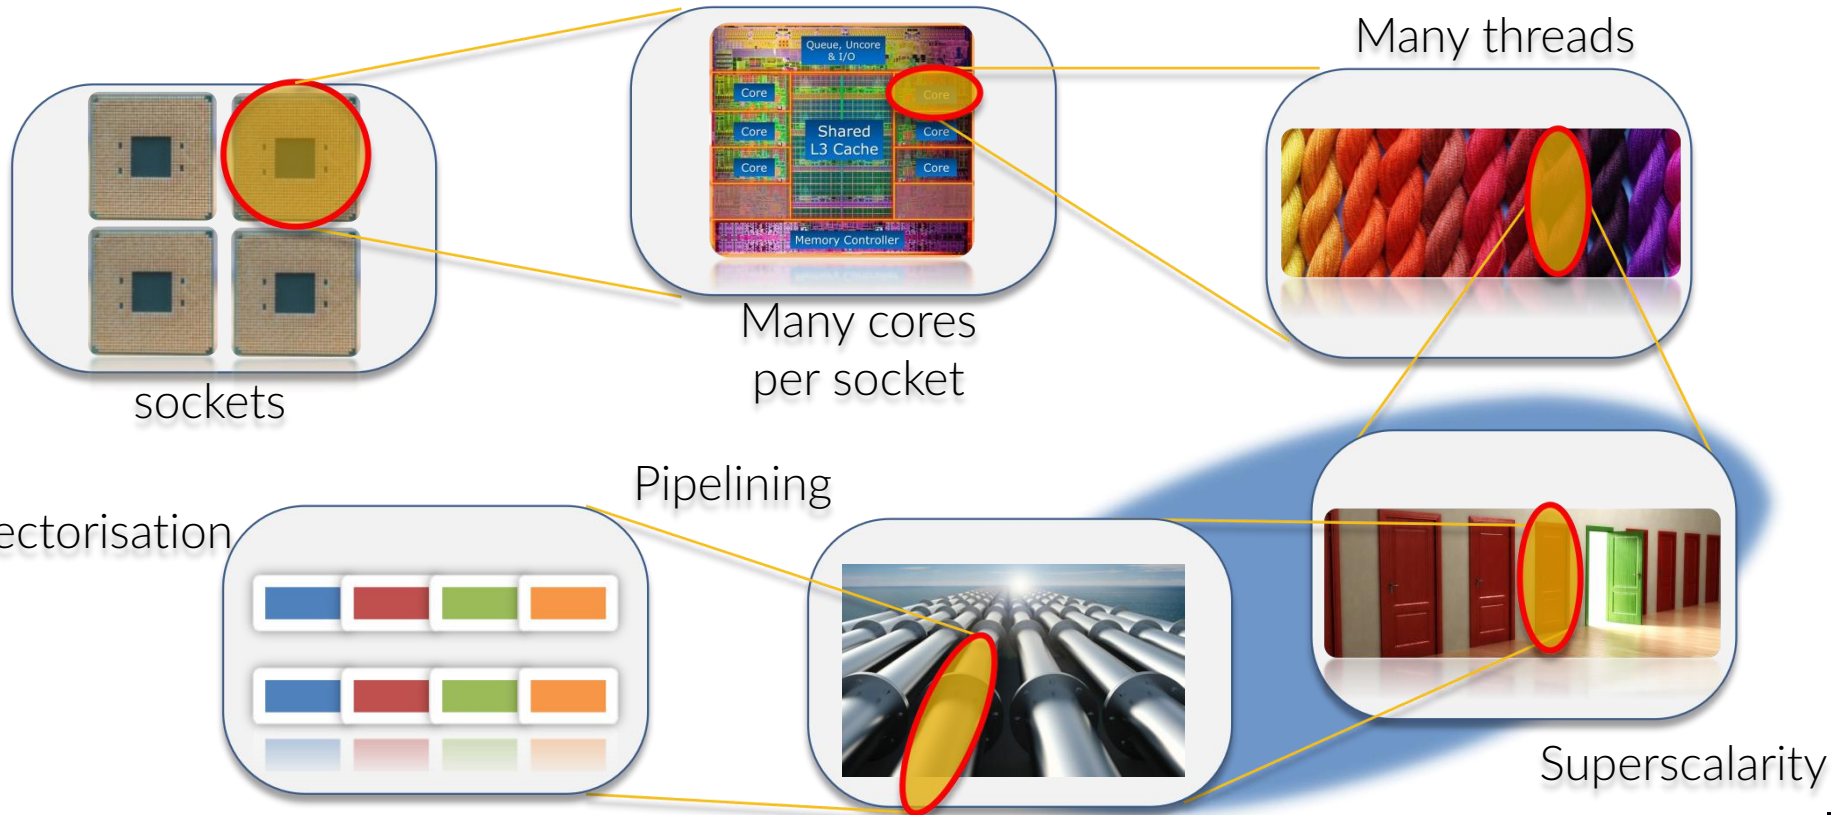


c





A final sketch



Vectorisation

that's all, have fun



“So long
and thanks
for all the fish”

A 3D rendered graphic featuring the text "So long and thanks for all the fish" in white, blocky, sans-serif letters. The text is arranged in three lines, with "So long" on the top line, "and thanks" in the middle, and "for all the fish" on the bottom line. The letters have a slight shadow and are set against a solid blue background. Above the text, there is a horizontal band with a dark wood grain texture. To the left and right of the text, there are small white 3D cubes that look like they are floating or falling.